

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



deeplearning.ai

Introduction to Named Entity Recognition

What is Named Entity Recognition?

- Locates and extracts predefined entities from text
- Places, organizations, names, time and dates



Types of Entities



Thailand:
Geographical



Google:
Organization



Indian:
Geopolitical

More Types of Entities



December:
Time Indicator




Egyptian statue:
Artifact



Barack Obama:
Person

Example of a labeled sentence



Applications of NER systems

- Search engine efficiency
- Recommendation engines
- Customer service
- Automatic trading






deeplearning.ai

Training NERs: Data Processing

Outline

- Convert words and entity classes into arrays
- Token padding
- Create a data generator



Processing data for NERs

- Assign each class a number
- Assign each word a number

Sharon flew to Miami last Friday.

[4282, 853, 187, 5388, 2894, 7]

B- O O B-geo O B-tim


per

Token padding

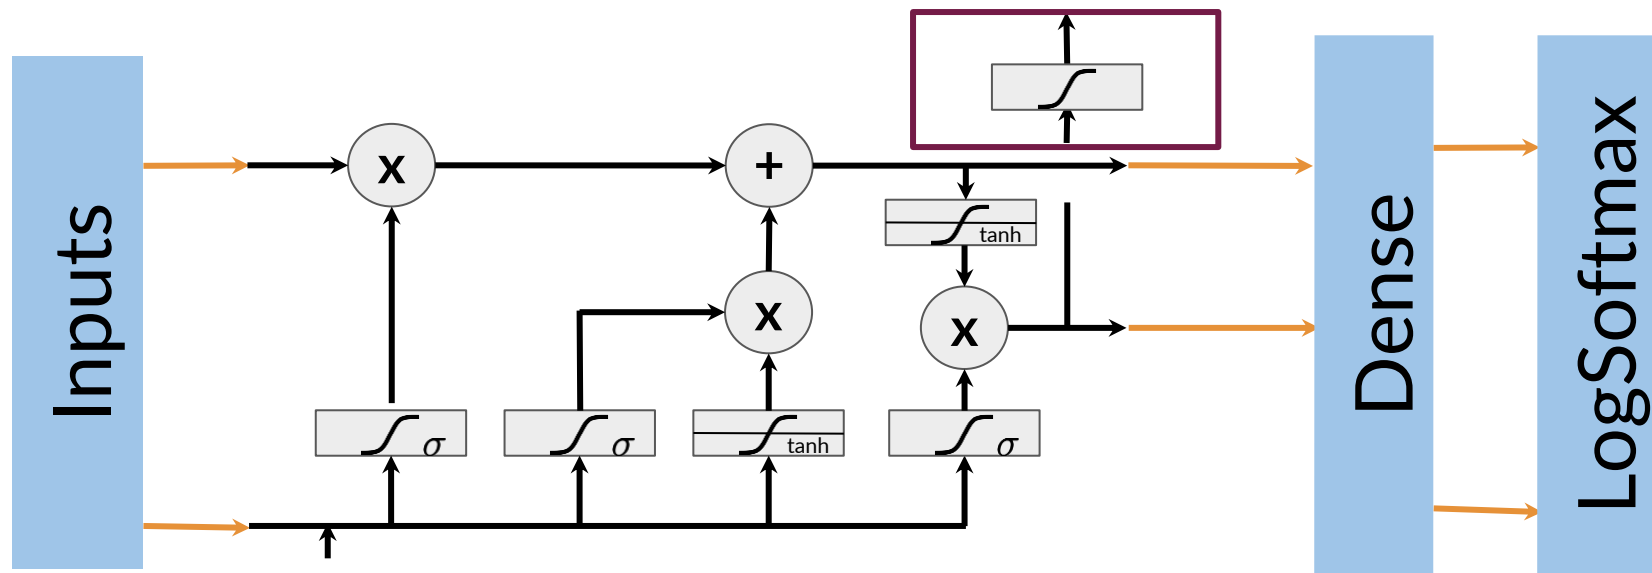
For LSTMs, all sequences need to be the same size.

- Set sequence length to a certain number
- Use the **<PAD>** token to fill empty spaces

Training the NER

1. Create a tensor for each input and its corresponding number
2. Put them in a batch  64, 128, 256, 512 ...
3. Feed it into an LSTM unit
4. Run the output through a dense layer
5. Predict using a log softmax over K classes

Training the NER



Layers in Trax

```
model = tl.Serial(  
    tl.Embedding(),  
    tl.LSTM(),  
    tl.Dense(),  
    tl.LogSoftmax()  
)
```

Summary

- Convert words and entities into same-length numerical arrays
- Train in batches for faster processing
- Run the output through a final layer and activation





deeplearning.ai

Computing Accuracy

Evaluating the model

1. Pass test set through the model
2. Get arg max across the prediction array
3. Mask padded tokens
4. Compare outputs against test labels

Evaluating the model in Python

```
def evaluate_model(test_sentences, test_labels, model):  
    pred = model(test_sentences)  
    outputs = np.argmax(pred, axis=2)  
    mask = ...  
    accuracy =   
    np.sum(outputs==test_labels)/float(np.sum(mask))  
  
    return accuracy
```

Summary

- If padding tokens, remember to mask them when computing accuracy
- Coding assignment!